

base_model_finetuning

April 11, 2025

```
[ ]: import tensorflow as tf
import tensorflow_datasets as tfds
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
```

```
[ ]: !wget https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/
↪extras/helper_functions.py
from helper_functions import plot_loss_curves, compare_histories
```

```
--2025-02-15 08:02:21-- https://raw.githubusercontent.com/mrdbourke/tensorflow-
deep-learning/main/extras/helper_functions.py
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 10246 (10K) [text/plain]
Saving to: 'helper_functions.py'
```

```
helper_functions.py 0%[                               ] 0 --.-KB/s
helper_functions.py 100%[=====>] 10.01K --.-KB/s in 0.001s
```

```
2025-02-15 08:02:21 (18.8 MB/s) - 'helper_functions.py' saved [10246/10246]
```

```
[ ]: # Set hyperparameters
BATCH_SIZE = 64
IMG_SIZE = 96 # Upscale CIFAR-10 images (32x32) to a larger size for ↪
↪MobileNetV2
AUTOTUNE = tf.data.AUTOTUNE
```

```
[ ]: (ds_train, ds_test), ds_info = tfds.load(
    'cifar10',
    split=['train', 'test'],
    as_supervised=True,
    with_info=True
)
```

```

Downloading and preparing dataset 162.17 MiB (download: 162.17 MiB, generated:
132.40 MiB, total: 294.58 MiB) to /root/tensorflow_datasets/cifar10/3.0.2...
Dl Completed...: 0 url [00:00, ? url/s]
Dl Size...: 0 MiB [00:00, ? MiB/s]
Extraction completed...: 0 file [00:00, ? file/s]
Generating splits...: 0%|          | 0/2 [00:00<?, ? splits/s]
Generating train examples...: 0%|          | 0/50000 [00:00<?, ? examples/s]
Shuffling /root/tensorflow_datasets/cifar10/incomplete.ICZA7D_3.0.2/
↳cifar10-train.tfrecord*...: 0%|          ...
Generating test examples...: 0%|          | 0/10000 [00:00<?, ? examples/s]
Shuffling /root/tensorflow_datasets/cifar10/incomplete.ICZA7D_3.0.2/cifar10-test.
↳tfrecord*...: 0%|          ...
Dataset cifar10 downloaded and prepared to
/root/tensorflow_datasets/cifar10/3.0.2. Subsequent calls will reuse this data.

```

```
[ ]: ds_info.features
```

```
[ ]: FeaturesDict({
    'id': Text(shape=(), dtype=string),
    'image': Image(shape=(32, 32, 3), dtype=uint8),
    'label': ClassLabel(shape=(), dtype=int64, num_classes=10),
})
```

```
[ ]: class_names = ds_info.features['label'].names
class_names
```

```
[ ]: ['airplane',
      'automobile',
      'bird',
      'cat',
      'deer',
      'dog',
      'frog',
      'horse',
      'ship',
      'truck']
```

```
[ ]: train_sample = ds_train.take(1) # takes one sample from train data
train_sample
```

```
[ ]: <_TakeDataset element_spec=(TensorSpec(shape=(32, 32, 3), dtype=tf.uint8,
name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>
```

```
[ ]: for image, label in train_sample:
    print(f'''
    image_shape : {image.shape}
    image_datatype : {image.dtype}
    target_class : {label}
    class_name : {class_names[label.numpy()]}
    ''')
    )
```

```
image_shape : (32, 32, 3)
image_datatype : <dtype: 'uint8'>
target_class : 7
class_name : horse
```

```
[ ]: import matplotlib.pyplot as plt
plt.figure(figsize=(1.5, 1.5))
plt.imshow(image)
plt.title(class_names[label.numpy()])
plt.axis(False)
```

```
[ ]: (-0.5, 31.5, 31.5, -0.5)
```

horse



```
[ ]: len(ds_train), len(ds_test)
```

```
[ ]: (50000, 10000)
```

```
[ ]: # data augmentation pipeline for training data
data_augmentation = tf.keras.Sequential([
    # Resize to desired IMG_SIZE
    tf.keras.layers.Resizing(IMG_SIZE, IMG_SIZE),
    tf.keras.layers.RandomFlip("horizontal"),
    tf.keras.layers.RandomRotation(0.1),
])
```

```
[ ]: # Preprocessing function using MobileNetV2's preprocess_input
def preprocess(image, label):
    image = tf.cast(image, tf.float32)
    image = preprocess_input(image)
    return image, label
```

```
[ ]: # Augmentation function (applied only on training set)
def augment(image, label):
    image = data_augmentation(image)
    return image, label
```

```
[ ]: # Prepare the training dataset
ds_train = ds_train.map(augment, num_parallel_calls=AUTOTUNE)
ds_train = ds_train.map(preprocess, num_parallel_calls=AUTOTUNE)
ds_train = ds_train.shuffle(1000).batch(BATCH_SIZE).prefetch(AUTOTUNE)
```

```
[ ]: # For the test dataset, only resize and preprocess (no augmentation)
def resize_and_preprocess(image, label):
    image = tf.cast(image, tf.float32)
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
    image = preprocess_input(image)
    return image, label
```

```
[ ]: ds_test = ds_test.map(resize_and_preprocess, num_parallel_calls=AUTOTUNE)
ds_test = ds_test.batch(BATCH_SIZE).prefetch(AUTOTUNE)
```

```
[ ]: # Build the model using MobileNetV2 as the base
base_model = MobileNetV2(
    include_top=False,
    weights='imagenet',
    input_shape=(IMG_SIZE, IMG_SIZE, 3)
)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_96_no_top.h5
9406464/9406464 0s
0us/step

```
[ ]: base_model.trainable = False
```

```
[ ]: inputs = tf.keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
x = base_model(inputs, training=False)
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.2)(x)
outputs = tf.keras.layers.Dense(10, activation='softmax')(x)
model = tf.keras.Model(inputs, outputs)
```

```
[ ]: model.compile(optimizer=tf.keras.optimizers.Adam(),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[ ]: model.summary()
```

Model: "functional_1"

Layer (type) ↳ Param #	Output Shape	
input_layer_2 (InputLayer) ↳ 0	(None, 96, 96, 3)	↳
mobilenetv2_1.00_96 (Functional) ↳ 2,257,984	(None, 3, 3, 1280)	↳
global_average_pooling2d ↳ 0 (GlobalAveragePooling2D) ↳	(None, 1280)	↳
dropout (Dropout) ↳ 0	(None, 1280)	↳
dense (Dense) ↳ 12,810	(None, 10)	↳

Total params: 2,270,794 (8.66 MB)

Trainable params: 12,810 (50.04 KB)


Non-trainable params: 2,257,984 (8.61 MB)


```
[ ]: history = model.fit(
    ds_train,
    epochs=5,
    validation_data=ds_test
)
```

Epoch 1/5
782/782 129s 149ms/step -

accuracy: 0.5983 - loss: 1.1976 - val_accuracy: 0.8204 - val_loss: 0.5330
Epoch 2/5
782/782 105s 132ms/step -
accuracy: 0.7455 - loss: 0.7420 - val_accuracy: 0.8322 - val_loss: 0.4826
Epoch 3/5
782/782 106s 133ms/step -
accuracy: 0.7520 - loss: 0.7080 - val_accuracy: 0.8306 - val_loss: 0.4913
Epoch 4/5
782/782 104s 131ms/step -
accuracy: 0.7584 - loss: 0.7036 - val_accuracy: 0.8342 - val_loss: 0.4910
Epoch 5/5
782/782 147s 138ms/step -
accuracy: 0.7661 - loss: 0.6819 - val_accuracy: 0.8390 - val_loss: 0.4653

```
[ ]: plot_loss_curves(history)
```





```
[ ]: base_model.trainable = True
fine_tune_at = 25
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False
```


```
[ ]: model.compile(optimizer=tf.keras.optimizers.Adam(1e-5),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```


```
[ ]: history = model.fit(
    ds_train,
    epochs=5,
    validation_data=ds_test,
)
```

```
Epoch 1/5
782/782          206s 197ms/step -
accuracy: 0.6337 - loss: 1.2428 - val_accuracy: 0.8408 - val_loss: 0.5099
Epoch 2/5
782/782          129s 163ms/step -
```

```
accuracy: 0.7762 - loss: 0.7042 - val_accuracy: 0.8672 - val_loss: 0.4018
Epoch 3/5
782/782          139s 160ms/step -
accuracy: 0.8082 - loss: 0.5913 - val_accuracy: 0.8781 - val_loss: 0.3795
Epoch 4/5
782/782          128s 162ms/step -
accuracy: 0.8330 - loss: 0.5092 - val_accuracy: 0.8806 - val_loss: 0.3760
Epoch 5/5
782/782          128s 161ms/step -
accuracy: 0.8473 - loss: 0.4601 - val_accuracy: 0.8815 - val_loss: 0.3751
```

```
[ ]: plot_loss_curves(history)
```





```
[ ]: history = model.fit(  
    ds_train,  
    epochs=3,  
    validation_data=ds_test,  
)
```

Epoch 1/3

782/782 130s 164ms/step -

accuracy: 0.8579 - loss: 0.4294 - val_accuracy: 0.8893 - val_loss: 0.3529

Epoch 2/3

782/782 142s 165ms/step -


accuracy: 0.8679 - loss: 0.3882 - val_accuracy: 0.8948 - val_loss: 0.3384


Epoch 3/3

782/782 131s 166ms/step -

accuracy: 0.8745 - loss: 0.3712 - val_accuracy: 0.8997 - val_loss: 0.3182

```
[ ]: plot_loss_curves(history)
```







```
[ ]: history = model.fit(  
    ds_train,  
    epochs=5,  
    validation_data=ds_test,  
)
```

```
Epoch 1/5  
782/782      133s 166ms/step -  
accuracy: 0.8846 - loss: 0.3405 - val_accuracy: 0.9032 - val_loss: 0.3093  
Epoch 2/5  
782/782      129s 163ms/step -  
accuracy: 0.8957 - loss: 0.3091 - val_accuracy: 0.9061 - val_loss: 0.3013  
Epoch 3/5  
782/782      128s 162ms/step -  
accuracy: 0.8972 - loss: 0.2990 - val_accuracy: 0.9092 - val_loss: 0.2870  
Epoch 4/5  
782/782      128s 161ms/step -  
accuracy: 0.9057 - loss: 0.2721 - val_accuracy: 0.9122 - val_loss: 0.2762  
Epoch 5/5  
782/782      125s 158ms/step -
```

accuracy: 0.9049 - loss: 0.2712 - val_accuracy: 0.9153 - val_loss: 0.2673

```
[ ]: plot_loss_curves(history)
```





```
[ ]: model.save("fine_tuned_model.keras")
      model.save("fine_tuned_model.h5", save_format='h5')
```

WARNING:absl:The `save_format` argument is deprecated in Keras 3. We recommend removing this argument as it can be inferred from the file path. Received: save_format=h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
[ ]: loaded_model = tf.keras.models.load_model("fine_tuned_model.keras")
```

```
[ ]: IMG_SIZE = 96  #matches the training size
      BATCH_SIZE = 64
      AUTOTUNE = tf.data.AUTOTUNE
```

```
[ ]: # Load full CIFAR-10 dataset
      ds_full, ds_info = tfds.load(
```

```
'cifar10',  
split='train+test',  
as_supervised=True,  
with_info=True  
)
```

```
[ ]: def preprocess(image, label):  
    image = tf.cast(image, tf.float32)  
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])  
    image = preprocess_input(image)  
    return image, label
```

```
[ ]: # Prepare dataset (apply preprocessing, batch, and prefetch)  
ds_full = ds_full.map(preprocess, num_parallel_calls=AUTOTUNE)  
ds_full = ds_full.batch(BATCH_SIZE).prefetch(AUTOTUNE)
```

```
[ ]: loss, accuracy = model.evaluate(ds_full)  
    loss, accuracy
```

```
938/938          27s 28ms/step -  
accuracy: 0.9364 - loss: 0.1922
```

```
[ ]: (0.16968880593776703, 0.9430000185966492)
```

```
[ ]:
```