

input_transformation

April 11, 2025

```
[ ]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
```

```
[ ]: # Hyperparameters
BATCH_SIZE = 64
IMG_SIZE = 96 # Upscale CIFAR-10 images (32x32) to 96x96 for MobileNetV2
AUTOTUNE = tf.data.AUTOTUNE
```

```
[ ]: def resize_and_preprocess(image, label):
    image = tf.cast(image, tf.float32)
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
    image = preprocess_input(image)
    return image, label
```

```
[ ]: # Load CIFAR-10 test dataset
(_, _), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
y_test = np.squeeze(y_test)
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071      13s
0us/step
```

```
[ ]: model = tf.keras.models.load_model("model.keras")
```

```
[ ]: #preprocessing data
test_dataset = tf.data.Dataset.from_tensor_slices((x_test, y_test))
test_dataset = test_dataset.map(resize_and_preprocess,
    ↪ num_parallel_calls=AUTOTUNE)
test_dataset = test_dataset.batch(BATCH_SIZE).prefetch(AUTOTUNE)
```

```
[ ]: loss, accuracy = model.evaluate(test_dataset)
```

```
157/157      13s 32ms/step -
accuracy: 0.9132 - loss: 0.2790
```

```
[ ]: loss, accuracy
```

```
[ ]: (0.2672818899154663, 0.9153000116348267)
```

```
[ ]: @tf.function
def batched_fgsm_attack(images, labels, epsilon=0.01):
    with tf.GradientTape() as tape:
        tape.watch(images)
        predictions = model(images, training=False)
        loss = tf.keras.losses.sparse_categorical_crossentropy(labels,
↪ predictions)
        gradients = tape.gradient(loss, images)
        adv_images = images + epsilon * tf.sign(gradients)
        adv_images = tf.clip_by_value(adv_images, -1, 1)
    return adv_images
```

```
[ ]: @tf.function
def batched_pgd_attack(images, labels, epsilon=0.01, alpha=0.005, num_iter=10):
    adv_images = tf.identity(images)

    for _ in tf.range(num_iter):
        with tf.GradientTape() as tape:
            tape.watch(adv_images)
            predictions = model(adv_images, training=False)
            loss = tf.keras.losses.sparse_categorical_crossentropy(labels,
↪ predictions)
            gradients = tape.gradient(loss, adv_images)
            adv_images = adv_images + alpha * tf.sign(gradients)

        # Project perturbation
        perturbation = tf.clip_by_value(adv_images - images, -epsilon, epsilon)
        adv_images = tf.clip_by_value(images + perturbation, -1, 1)

    return adv_images
```

```
[ ]: def deepfool_attack(image, num_classes=10, overshoot=0.0000001, max_iter=1):
    image = tf.convert_to_tensor(image, dtype=tf.float32)
    perturbed_image = tf.identity(image)

    # Get original prediction and label
    with tf.GradientTape() as tape:
        tape.watch(perturbed_image)
        logits = model(tf.expand_dims(perturbed_image, axis=0))[0]
    orig_label = tf.argmax(logits)

    r_tot = tf.zeros_like(image)
    i = 0

    while i < max_iter:
```

```

with tf.GradientTape(persistent=True) as tape:
    tape.watch(perturbed_image)
    logits = model(tf.expand_dims(perturbed_image, axis=0))[0]

current_label = tf.argmax(logits)
if current_label != orig_label:
    break

# Compute gradients for all class logits
gradients = []
for k in range(num_classes):
    with tf.GradientTape() as tape2:
        tape2.watch(perturbed_image)
        logit_k = model(tf.expand_dims(perturbed_image, axis=0))[0, k]
        grad_k = tape2.gradient(logit_k, perturbed_image)
        gradients.append(grad_k)
gradients = tf.stack(gradients)

# Compute minimal perturbation
f_orig = logits[orig_label]
perturbs = []
for k in range(num_classes):
    if k == orig_label:
        continue
    w_k = gradients[k] - gradients[orig_label]
    f_k = logits[k] - f_orig
    norm_w = tf.norm(tf.reshape(w_k, [-1])) + 1e-8
    pert_k = tf.abs(f_k) / norm_w
    perturbs.append((pert_k, w_k))

# Choose the closest decision boundary
perturbs.sort(key=lambda x: x[0])
pert_k, w_k = perturbs[0]

# Compute minimal directional perturbation (no sign scaling)
r_i = (pert_k * w_k) / (tf.norm(w_k) + 1e-8)
r_tot += r_i

# Apply accumulated perturbation with small overshoot
perturbed_image = image + (1 + overshoot) * r_tot
perturbed_image = tf.clip_by_value(perturbed_image, -1, 1)

i += 1

return perturbed_image

```

```
[ ]: def get_test_dataset():
    # Load CIFAR-10 test dataset and preprocess
    (_, _), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
    y_test = np.squeeze(y_test)
    ds = tf.data.Dataset.from_tensor_slices((x_test, y_test))
    ds = ds.map(resize_and_preprocess, num_parallel_calls=AUTOTUNE)
    ds = ds.batch(BATCH_SIZE).prefetch(AUTOTUNE)
    return ds
```

```
[ ]: clean_ds = get_test_dataset()
model.compile(loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
[ ]: def build_adversarial_dataset_fast(dataset, attack_fn, attack_name="FGSM"):
    adv_images_all = []
    adv_labels_all = []

    print(f"\nBuilding {attack_name} dataset...")

    for images, labels in dataset:
        adv_images = attack_fn(images, labels)
        adv_images_all.append(adv_images)
        adv_labels_all.append(labels)

    adv_images_all = tf.concat(adv_images_all, axis=0)
    adv_labels_all = tf.concat(adv_labels_all, axis=0)

    adv_ds = tf.data.Dataset.from_tensor_slices((adv_images_all,
    ↪adv_labels_all))
    return adv_ds.batch(BATCH_SIZE).prefetch(AUTOTUNE)
```

```
[ ]: def build_adversarial_dataset_deepfool(attack_fn, name="DeepFool",
    ↪max_samples=500, num_classes=10):
    adv_images = []
    adv_labels = []

    print(f"\nGenerating {name} adversarial dataset (max {max_samples} samples).
    ↪..")
    sample_count = 0

    for images, labels in clean_ds:
        for img, label in zip(images, labels):
            # Pass a fixed number of classes instead of the label value.
            adv_img = attack_fn(img, num_classes)
            adv_images.append(adv_img.numpy())
            adv_labels.append(int(label.numpy()))
            sample_count += 1
```

```

        if sample_count >= max_samples:
            break
    if sample_count >= max_samples:
        break

    adv_images = np.array(adv_images)
    adv_labels = np.array(adv_labels)

    ds = tf.data.Dataset.from_tensor_slices((adv_images, adv_labels))
    ds = ds.batch(BATCH_SIZE).prefetch(AUTOTUNE)
    return ds

```

```

[ ]: def evaluate_model_on_dataset(dataset, name="Dataset"):
    y_true, y_pred = [], []
    total_loss = 0.0
    total_samples = 0

    loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
    for batch_images, batch_labels in dataset:
        preds = model(batch_images, training=False)
        loss = loss_fn(batch_labels, preds).numpy()
        pred_classes = tf.argmax(preds, axis=1).numpy()
        y_true.extend(batch_labels.numpy())
        y_pred.extend(pred_classes)
        total_loss += loss * len(batch_labels)
        total_samples += len(batch_labels)

    accuracy = np.mean(np.array(y_true) == np.array(y_pred))
    avg_loss = total_loss / total_samples
    correct = sum(np.array(y_true) == np.array(y_pred))
    incorrect = total_samples - correct

    print(f"\n{name} Evaluation:")
    print(f"  Total Samples: {total_samples}")
    print(f"  Accuracy: {accuracy:.4f}")
    print(f"  Loss: {avg_loss:.4f}")
    print(f"  Correct Predictions: {correct}")
    print(f"  Incorrect Predictions: {incorrect}")
    return accuracy, avg_loss

```

```

[ ]: fgsm_ds = build_adversarial_dataset_fast(clean_ds, lambda x, y:
    ↳ batched_fgsm_attack(x, y, epsilon=0.01), attack_name="FGSM")
pgd_ds = build_adversarial_dataset_fast(clean_ds, lambda x, y:
    ↳ batched_pgd_attack(x, y, epsilon=0.01, alpha=0.005, num_iter=10),
    ↳ attack_name="PGD")

```

Building FGSM dataset...

Building PGD dataset...

```
[ ]: evaluate_model_on_dataset(fgsm_ds, name="FGSM")
```

FGSM Evaluation:

Total Samples: 10000

Accuracy: 0.1820

Loss: 5.1916

Correct Predictions: 1820

Incorrect Predictions: 8180

```
[ ]: (np.float64(0.182), np.float32(5.191604))
```

```
[ ]: evaluate_model_on_dataset(pgd_ds, name="PGD")
```

PGD Evaluation:

Total Samples: 10000

Accuracy: 0.0000

Loss: 22.0664

Correct Predictions: 0

Incorrect Predictions: 10000

```
[ ]: (np.float64(0.0), np.float32(22.066427))
```

```
[ ]: deepfool_ds = build_adversarial_dataset_deepfool(deepfool_attack,
↳ name="DeepFool", max_samples=200)
evaluate_model_on_dataset(deepfool_ds, name="DeepFool Attack")
```

Generating DeepFool adversarial dataset (max 200 samples)...

DeepFool Attack Evaluation:

Total Samples: 200

Accuracy: 0.1500

Loss: 5.1534

Correct Predictions: 30

Incorrect Predictions: 170

```
[ ]: (np.float64(0.15), np.float32(5.153436))
```

```
[ ]: def get_gaussian_kernel(size=3, sigma=1.0):
    """Creates a 2D Gaussian kernel."""
    x = tf.range(-size // 2 + 1, size // 2 + 1, dtype=tf.float32)
    x = tf.exp(-(x**2) / (2 * sigma**2))
```

```

kernel_1d = x / tf.reduce_sum(x)
kernel_2d = tf.tensordot(kernel_1d, kernel_1d, axes=0)
kernel_2d = kernel_2d / tf.reduce_sum(kernel_2d)
return kernel_2d[:, :, tf.newaxis, tf.newaxis] # Shape: (H, W, 1, 1)
↳ in_channels=1, out_channels=1)

def apply_gaussian_blur(x, sigma):
    """Applies Gaussian blur using depthwise convolution."""
    kernel = get_gaussian_kernel(size=3, sigma=sigma)
    channels = tf.shape(x)[-1]
    kernel = tf.tile(kernel, [1, 1, channels, 1]) # Make kernel channel-wise
    x = tf.nn.depthwise_conv2d(x, kernel, strides=[1, 1, 1, 1], padding='SAME')
    return x

```

```

[ ]: def inference_input_transformation(
    x,
    apply_bitdepth=True,
    bits=4,
    apply_noise=True,
    noise_std=0.05,
    apply_jpeg=True,
    jpeg_quality=75,
    apply_blur=True,
    blur_sigma=0.5
):
    """
    Apply input transformations: quantization, noise, JPEG compression, and
    ↳ blur.

    Args:
        x (Tensor): Input tensor in [0,1].
    """
    if apply_bitdepth:
        levels = 2 ** bits
        x = tf.round(x * (levels - 1)) / (levels - 1)

    if apply_noise:
        noise = tf.random.normal(tf.shape(x), mean=0.0, stddev=noise_std,
        ↳ dtype=x.dtype)
        x = x + noise

    if apply_jpeg:
        def jpeg_fn(img):
            img_uint8 = tf.image.convert_image_dtype(img, tf.uint8)
            encoded = tf.io.encode_jpeg(img_uint8, quality=jpeg_quality)
            decoded = tf.io.decode_jpeg(encoded)
            return tf.image.convert_image_dtype(decoded, tf.float32)

```

```

        x = tf.map_fn(jpeg_fn, x)

    if apply_blur:
        x = apply_gaussian_blur(x, sigma=blur_sigma)

    x = tf.clip_by_value(x, 0.0, 1.0)
    return x

```

```

[ ]: class TransformedModel(tf.keras.Model):
    def __init__(
        self,
        base_model,
        apply_bitdepth=True,
        bits=4,
        apply_noise=True,
        noise_std=0.05,
        apply_jpeg=True,
        jpeg_quality=75,
        apply_blur=True,
        blur_sigma=0.5
    ):
        super().__init__()
        self.base_model = base_model
        self.apply_bitdepth = apply_bitdepth
        self.bits = bits
        self.apply_noise = apply_noise
        self.noise_std = noise_std
        self.apply_jpeg = apply_jpeg
        self.jpeg_quality = jpeg_quality
        self.apply_blur = apply_blur
        self.blur_sigma = blur_sigma

    def call(self, inputs, training=False):
        # Convert from [-1, 1] to [0, 1] before transformation
        inputs = (inputs + 1.0) / 2.0

        transformed = inference_input_transformation(
            inputs,
            apply_bitdepth=self.apply_bitdepth,
            bits=self.bits,
            apply_noise=self.apply_noise,
            noise_std=self.noise_std,
            apply_jpeg=self.apply_jpeg,
            jpeg_quality=self.jpeg_quality,
            apply_blur=self.apply_blur,
            blur_sigma=self.blur_sigma
        )

```



```
# Convert back to [-1, 1] for model input
transformed = transformed * 2.0 - 1.0
return self.base_model(transformed, training=training)
```

```
[ ]: model = TransformedModel(model)
```

```
[ ]: evaluate_model_on_dataset(clean_ds, name='Clean + transformed')
```

Clean + transformed Evaluation:

```
Total Samples: 10000
Accuracy: 0.4095
Loss: 2.4089
Correct Predictions: 4095
Incorrect Predictions: 5905
```

```
[ ]: (np.float64(0.4095), np.float32(2.4088583))
```

```
[ ]: evaluate_model_on_dataset(fgsm_ds, name="FGSM + Transformed")
```

FGSM + Transformed Evaluation:

```
Total Samples: 10000
Accuracy: 0.3898
Loss: 2.5223
Correct Predictions: 3898
Incorrect Predictions: 6102
```

```
[ ]: (np.float64(0.3898), np.float32(2.5223048))
```

```
[ ]: evaluate_model_on_dataset(pgd_ds, name="PGD + Transformed")
```

PGD + Transformed Evaluation:

```
Total Samples: 10000
Accuracy: 0.3863
Loss: 2.4916
Correct Predictions: 3863
Incorrect Predictions: 6137
```

```
[ ]: (np.float64(0.3863), np.float32(2.4916496))
```

```
[ ]: evaluate_model_on_dataset(deepfool_ds, name="DeepFool + Transformed")
```

DeepFool + Transformed Evaluation:

```
Total Samples: 200
```

```
Accuracy: 0.1900  
Loss: 4.7012  
Correct Predictions: 38  
Incorrect Predictions: 162
```

```
[ ]: (np.float64(0.19), np.float32(4.7012343))
```

```
[ ]:
```